

The initial step in designing any model is to understand the data, including its shape, values, and any categorical features. When we first encounter the data, it appears in the following way.

	Elevation	Aspect	Slope	Horizontal_Distance_To_Hydrology	\
0	2596	51	3		258
1	2590	56	2		212
2	2804	139	9		268
3	2785	155	18		242
4	2595	45	2		153

	Vertical_Distance_To_Hydrology	Horizontal_Distance_To_Roadways	\
0	0	510	
1	-6	398	
2	65	3100	
3	118	3090	
4	-1	391	

	Hillshade_9am	Hillshade_Noon	Hillshade_3pm	\
0	221	232	148	
1	220	235	151	
2	234	238	135	
3	238	238	122	
4	220	234	150	

	Horizontal_Distance_To_Fire_Points	...	Soil_Type32	Soil_Type33	\
0	6279	...	0	0	
1	6225	...	0	0	
2	6121	...	0	0	
3	6211	...	0	0	
4	6172	...	0	0	

	Soil_Type34	Soil_Type35	Soil_Type36	Soil_Type37	Soil_Type38	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	Soil_Type39	Soil_Type40	Cover_Type	\
0	0	0	5	
1	0	0	5	
2	0	0	2	
3	0	0	2	
4	0	0	5	

[5 rows x 55 columns]

Our dataset consists of 55 columns, representing 54 features for each item, with the last column being our target variable. It is apparent that our data does not contain any categorical features. However, it is worth noting that certain columns have significantly larger values compared to others, indicating the need for scaling in future steps.

After dividing our data into train and test sets, we examined them for null and duplicate values, but found none. For preprocessing, we only applied scaling to our first column, which produced the following results.

	Elevation	Aspect	Slope	Horizontal_Distance_To_Hydrology	\
183380	0.702466	-0.784220	0.386543		-0.683720
561389	-1.431609	1.772207	-0.681631		-1.069421
370958	0.003004	-0.962991	-0.414587		3.126252
429888	-0.264648	-1.231147	-0.815152		-1.125865
178657	-0.821363	-0.650141	1.187673		-0.561425

	Vertical_Distance_To_Hydrology	Horizontal_Distance_To_Roadways	\
183380	-0.315721	1.946292	
561389	-0.779104	-0.152020	
370958	2.533231	0.192671	
429888	-0.779104	-0.348435	
178657	-0.092610	-1.409466	

	Hillshade_9am	Hillshade_Noon	Hillshade_3pm	\
183380	0.853218	-0.977752	-1.241733	
561389	-0.265227	0.034056	0.376652	
370958	0.443122	-0.421257	-0.510850	
429888	0.107588	0.034056	0.089519	
178657	1.188752	-1.483656	-1.946513	

	Horizontal_Distance_To_Fire_Points	...	Soil_Type31	Soil_Type32	\
183380	0.308453	...	0	0	
561389	0.170209	...	0	0	
370958	-0.947828	...	1	0	
429888	-0.009583	...	0	0	
178657	0.376442	...	0	0	

	Soil_Type33	Soil_Type34	Soil_Type35	Soil_Type36	Soil_Type37	\
183380	0	0	0	0	0	
561389	0	0	0	0	0	
370958	0	0	0	0	0	

The columns are now more uniform in their values, as can be observed. It's worth noting that the varying IDs are due to the shuffling that occurred during the data splitting process.

The different parameters we used for the decision tree models were 'max\_depth': [20, 25], 'criterion': ['entropy', 'log\_loss'], 'min\_samples\_split': [5, 7], and 'min\_samples\_leaf': [1, 2], resulting in a total of 16 models. After testing various combinations, we found that leaving the max depth at its default value resulted in a better model. Additionally, we discovered that the gini criterion typically performed worse for this particular dataset, so we excluded it from our analysis. The results of our analysis are presented below.

Best parameters for Decision Tree: {'criterion': 'log\_loss', 'max\_depth': 25, 'min\_samples\_leaf': 1, 'min\_samples\_split': 5}

```


params \
12 {'criterion': 'log_loss', 'max_depth': 25, 'min_samples_leaf': 1, 'min_samples_split': 5}
4  {'criterion': 'entropy', 'max_depth': 25, 'min_samples_leaf': 1, 'min_samples_split': 5}
13 {'criterion': 'log_loss', 'max_depth': 25, 'min_samples_leaf': 1, 'min_samples_split': 7}
5  {'criterion': 'entropy', 'max_depth': 25, 'min_samples_leaf': 1, 'min_samples_split': 7}
6  {'criterion': 'entropy', 'max_depth': 25, 'min_samples_leaf': 2, 'min_samples_split': 5}
14 {'criterion': 'log_loss', 'max_depth': 25, 'min_samples_leaf': 2, 'min_samples_split': 5}
7  {'criterion': 'entropy', 'max_depth': 25, 'min_samples_leaf': 2, 'min_samples_split': 7}
15 {'criterion': 'log_loss', 'max_depth': 25, 'min_samples_leaf': 2, 'min_samples_split': 7}
0  {'criterion': 'entropy', 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 5}
8  {'criterion': 'log_loss', 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 5}
1  {'criterion': 'entropy', 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 7}
9  {'criterion': 'log_loss', 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 7}
10 {'criterion': 'log_loss', 'max_depth': 20, 'min_samples_leaf': 2, 'min_samples_split': 5}
2  {'criterion': 'entropy', 'max_depth': 20, 'min_samples_leaf': 2, 'min_samples_split': 5}
3  {'criterion': 'entropy', 'max_depth': 20, 'min_samples_leaf': 2, 'min_samples_split': 7}
11 {'criterion': 'log_loss', 'max_depth': 20, 'min_samples_leaf': 2, 'min_samples_split': 7}

```

	mean_test_score	std_test_score
12	0.925	0.002
4	0.925	0.002
13	0.924	0.002
5	0.923	0.002
6	0.923	0.002
14	0.923	0.002
7	0.922	0.002
15	0.921	0.002
0	0.904	0.003
8	0.904	0.003
1	0.903	0.003
9	0.903	0.003
10	0.903	0.003
2	0.903	0.003
3	0.902	0.003
11	0.902	0.003
done		

The optimal parameter set for decision tree is {'criterion': 'log\_loss', 'min\_samples\_leaf': 1, 'min\_samples\_split': 5}. The result implies that log\_loss and entropy typically perform similarly, so we used log\_loss for the parameters in our random forest.

The parameter range we considered for random forest was {'n\_estimators': [100, 150], 'criterion': ['entropy'], 'min\_samples\_split': [5, 7], 'min\_samples\_leaf': [1, 2]}.



Which resulted:


```
-----
Best parameters for Random Forest: {'criterion': 'entropy', 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 150}

params \
1 {'criterion': 'entropy', 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 150}
0 {'criterion': 'entropy', 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}
3 {'criterion': 'entropy', 'min_samples_leaf': 1, 'min_samples_split': 7, 'n_estimators': 150}
2 {'criterion': 'entropy', 'min_samples_leaf': 1, 'min_samples_split': 7, 'n_estimators': 100}
5 {'criterion': 'entropy', 'min_samples_leaf': 2, 'min_samples_split': 5, 'n_estimators': 150}
4 {'criterion': 'entropy', 'min_samples_leaf': 2, 'min_samples_split': 5, 'n_estimators': 100}
7 {'criterion': 'entropy', 'min_samples_leaf': 2, 'min_samples_split': 7, 'n_estimators': 150}
6 {'criterion': 'entropy', 'min_samples_leaf': 2, 'min_samples_split': 7, 'n_estimators': 100}

mean_test_score  std_test_score
1      0.948      8.080e-04
0      0.947      9.276e-04 |
3      0.945      8.421e-04
2      0.945      8.932e-04
5      0.944      9.628e-04
4      0.944      7.413e-04
7      0.942      1.302e-03
6      0.941      1.284e-03
```

The accuracy obtained by the random forest model was better, which is in line with what we were hoping to achieve.

The optimal set of parameters for this model were {'criterion': 'entropy', 'min\_samples\_leaf': 1, 'min\_samples\_split': 5, 'n\_estimators': 150}.



The X shows the combination of parameters used, and the values are represented in logarithmic form to highlight the differences.

After fitting our model, it is time to evaluate it using different metrics. Let's begin by describing the metrics we will be discussing:

- **Accuracy:** Accuracy measures the overall correctness of the model's predictions. It is calculated as the number of correct predictions divided by the total number of predictions.
- **Precision:** Precision measures how many of the predicted positive cases were actually positive. It is calculated as the number of true positives divided by the sum of true positives and false positives.
- **Recall:** Recall measures how many of the actual positive cases were correctly predicted as positive by the model. It is calculated as the number of true positives divided by the sum of true positives and false negatives.
- **F1 score:** F1 score is a weighted average of precision and recall, and it balances between the two metrics. It is calculated as 2 times the product of precision and recall, divided by the sum of precision and recall.

Confusion matrix: which in our case is 7\*7 because it's not a binary classification.

	PREDICTED CLASS	
	Class=Yes	Class=No
	Class=Yes	Class=No
ACTUAL CLASS	a	b
	c	d

a: TP (true positive)

b: FN (false negative)

c: FP (false positive)

d: TN (true negative)

Now we should analyze this metrics for our predicted values.

Decision Tree:

```
0.9391694969708096
      precision    recall  f1-score   support

     1       0.94       0.94       0.94      63756
     2       0.95       0.95       0.95      84864
     3       0.93       0.94       0.93     10718
     4       0.85       0.83       0.84       782
     5       0.84       0.82       0.83      2875
     6       0.89       0.88       0.88      5182
     7       0.95       0.94       0.95      6127

 accuracy      0.94      174304
 macro avg     0.91      174304
 weighted avg  0.94      174304

[[59969 3482   3    0   51   7  244]
 [ 3730 80362 232   0  346 145   49]
 [    3   192 10049 79   28  367   0]
 [    0    0   95  652   0   35   0]
 [   52  423   30   0 2357  12   1]
 [    8   173  400  36   9 4556   0]
 [  335   34    1   0   1   0 5756]]
```

As you can see, our accuracy is 94%. For the other metrics, let's consider the 4th class. Its precision is 85%, which means 85% of the items we predicted as the 4th class truly belong to the 4th class. Its recall is 83%, which means 83% of the 4th class items were predicted correctly. Finally, the F1 measure can be computed as  $(2 \times 0.85 \times 0.83) / (0.85 + 0.83)$ .

Other classes can be computed in the same way using the confusion matrix. All the data on the diagonal are the ones that were predicted correctly, and precision is calculated as the data on the diagonal divided by the sum of its column. Recall is also computed by dividing the desired diagonal data by its row.

## Random Forest:

```

-----
              precision    recall  f1-score   support

         1         0.97         0.94         0.95         63756
         2         0.95         0.97         0.96         84864
         3         0.94         0.96         0.95         10718
         4         0.91         0.86         0.89           782
         5         0.94         0.77         0.85         2875
         6         0.93         0.89         0.91         5182
         7         0.97         0.94         0.96         6127

 accuracy                   0.95         174304
 macro avg         0.94         0.91         0.92         174304
 weighted avg      0.95         0.95         0.95         174304

[[59901 3704    1    0    14    4   132]
 [ 1774 82664  182    1   112   103   28]
 [    2   152 10294   40    16   214    0]
 [    0    0    90  674    0    18    0]
 [   37   554   41    0  2227   16    0]
 [    6   163  362   23    3  4625    0]
 [   320   24    0    0    1    0  5782]]

```

The accuracy of our model is 0.95%. For the precision of the 4th class, we found that 0.91% of the items we predicted as the 4th class are actually a part of the 4th class. The recall of the 4th class is 0.96, meaning that 0.96% of the 4th class items were predicted correctly. To calculate the F1 score, we can use the formula  $(2 \times 0.91 \times 0.96) / (0.96 + 0.91)$ . We can use the confusion matrix to calculate these metrics for other classes as well. The diagonal entries in the matrix represent the items that were predicted correctly, and precision is calculated by dividing the diagonal entry by the sum of its column. Recall is calculated by dividing the diagonal entry by the sum of its row.